

C28x Compiler Optimization

George Mock

August 2019

Agenda

- Enable Optimization
- Prefer Signed Loop Counters
- Background
- Restrict
- Dual MAC – C Code Method
- Dual MAC – Intrinsic Method

Enable Optimization

- Focus on inner loop
- How many assembly instructions?
- --opt_level=off
- --opt_level=2

```
void f_auto_correlation(AUTO_CORR *f)
{
    int i, k;
    float sum;

    for (i=0; i < f->lag; i++)
    {
        sum = 0;

        /* ----- INNER LOOP ----- */
        for (k = f->lag; k < (f->len+f->lag); k++)
            sum += f->input[k] * f->input[k-i];

        f->output[i] = sum;
    }
}
```

Inner Loop Comparison

--opt_level=off (default)

```
$C$L2:
    MOV     ACC,*-SP[6] << 1
    ADDL    ACC,*+XAR5[2]
    MOVL    XAR4,ACC
    MOV     AL,*-SP[6]
    SUB     AL,*-SP[5]
    MOV32   R0H,*+XAR4[0]
    MOV     ACC,AL << 1
    ADDL    ACC,*+XAR5[2]
    MOVL    XAR4,ACC
    MOV32   R1H,*+XAR4[0]
    MPYF32  R1H,R1H,R0H
    MOV32   R3H,*-SP[4]
    ADDF32  R0H,R1H,R3H
    NOP
    MOV32   *-SP[4],R0H
    INC     *-SP[6]
$C$L3:
    MOVL    XAR5,*-SP[2]
    MOVL    XAR4,*-SP[2]
    MOV     AL,*+XAR5[1]
    ADD     AL,*+XAR4[0]
    CMP     AL,*-SP[6]
    B       $C$L2,GT
```

--opt_level=2

```

RPT     AR5
|| MACF32  R7H,R3H,*XAR6++,*XAR7++
    ADDF32  R3H,R3H,R2H
    ADDF32  R2H,R7H,R6H
    NOP
    ADDF32  R3H,R3H,R2H
```

Agenda

- Enable Optimization
- **Prefer Signed Loop Counters**
- Background
- Restrict
- Dual MAC – C Code Method
- Dual MAC – Intrinsic Method

Prefer Signed Loop Counters

```
int i, j;    /* LOOP COUNTERS */
float sum;
for (i=0; i < N; i++)
    for (j=0; j < M; j++)
        ...
```

- Prefer signed over unsigned
 - `int` instead of `unsigned int`
- Compiler must presume unsigned counters may *wrap around* from 0xffff to 0
- Compiler assumes signed counters never wrap
 - If they do wrap, that is user error
- When `i` is signed, the compiler often changes `array[i]` into `*XAR3++`

Agenda

- Enable Optimization
- Prefer Signed Loop Counters
- **Background**
- Restrict
- Dual MAC – C Code Method
- Dual MAC – Intrinsic Method

Do NOT Lie to the Compiler

- Many methods for giving the compiler extra information
 - restrict
 - #pragma
 - intrinsic
- Verify the extra information is always correct
 - Under every circumstance
 - For every call
- Document constraints to those who call your functions
- Lying can cause bugs that are very hard to find
 - No diagnostics
 - Program silently does the wrong thing!

What is an Intrinsic?

- Similar to a function call
- Not implemented by calling a function
- Instead, a few instructions are emitted
 - Often, just one instruction
- An effective way to access unusual HW features of C28x
- Examples

```
long_var1 = __rol(long_var2);          /* rotate left */  
long_var1 = __sat(long_var2);          /* saturate    */  
__dmac(p1_32[i], p2_32[i], acc1, acc2, 0); /* dual MAC    */
```

Agenda

- Enable Optimization
- Prefer Signed Loop Counters
- Background
- **Restrict**
- Dual MAC – C Code Method
- Dual MAC – Intrinsic Method

Restrict Example

```
typedef struct {  
    float *p1;  
    float *p2;  
    float prod, G;  
    int length;  
} VECTOR;  
  
void vector_dot_product_squared(VECTOR *vec)  
{  
    int i;  
    for (i = 0; i < vec->length; i++)  
    {  
        vec->prod += vec->p1[i] * vec->p2[i];  
        vec->G     += vec->p2[i] * vec->p2[i];  
    }  
}
```

Restrict Details

```
typedef struct {  
    float * restrict p1;  
    float * restrict p2;  
    float prod, G;  
    int length;  
} VECTOR;
```

- Put **restrict** between * and the name of the pointer variable
- Property of the pointer, not the memory locations accessed through the pointer
- Data accessed through a restrict pointer is never accessed another way
 - During the scope of the pointer
 - Conservatively correct definition
 - Full definition allows corner cases that rarely matter in practice

Restrict Example

- Could `p1` point to `prod` or `G`?
- Could `p2` point to `prod` or `G`?
- Default answer is yes
- Restrict changes the answer to no
- What effect does that have?

```
for (i = 0; i < vec->length; i++)  
{  
    vec->prod += vec->p1[i] * vec->p2[i];  
    vec->G    += vec->p2[i] * vec->p2[i];  
}
```

Restrict Comparison

No restrict (default)

```

; repeat block starts      ; loop starts
$C$L1:
    MOV32    R1H,*XAR7++    ; read
    MOV32    R2H,*+XAR5[0]  ; read
    MPYF32   R2H,R2H,R1H
    NOP
    ADDF32   R0H,R0H,R2H
    NOP
    MOV32    *+XAR4[4],R0H   ; write
    MOV32    R1H,*XAR5++    ; read
    MPYF32   R2H,R1H,R1H
    NOP
    ADDF32   R3H,R3H,R2H
    NOP
    MOV32    *+XAR4[6],R3H   ; write
    ; repeat block ends    ; loop ends
$C$L2:
    LRETR
  
```

Add restrict

```

; repeat block starts      ; loop starts
$C$L1:
    MOV32    R2H,*XAR7++    ; read

    MPYF32   R1H,R2H,R2H
    ||      MOV32    R4H,*XAR5++    ; read

    MPYF32   R2H,R2H,R4H
    ADDF32   R0H,R0H,R1H
    ADDF32   R3H,R3H,R2H
    NOP
    ; repeat block ends    ; loop ends
$C$L2:
    MOV32    *+XAR4[6],R0H   ; write
    MOV32    *+XAR4[4],R3H   ; write
$C$L3:
    MOV32    R4H,*--SP
    LRETR
  
```

Agenda

- Enable Optimization
- Prefer Signed Loop Counters
- Background
- Restrict
- **Dual MAC – C Code Method**
- Dual MAC – Intrinsic Method

Dual MAC

- DMAC is most powerful instruction on C28x
- How can it be generated from C code?
- Two methods
 - C code with extra information
 - Intrinsic

Dual Mac – C Code Method

```
#include <stdint.h>

int32_t compute_mac(int16_t *p1,
                    int16_t *p2, int_fast16_t length)
{
    int_fast16_t i;
    int32_t result = 0;

    _nassert((intptr_t)p1 % 2 == 0);
    _nassert((intptr_t)p2 % 2 == 0);

    #pragma MUST_ITERATE(, , 2)
    for (i = 0; i < length; i++)
        result += (int32_t) p1[i] * p2[i];

    return result;
}
```

- Method: C code with extra information
- The next few slides repeat key lines from this example, then explain them

Include <stdint.h>

```
#include <stdint.h>
```

- Use standardized type names from <stdint.h>

Type	Means
int32_t	signed, exactly 32-bits
int16_t	signed, exactly 16-bits
int_fast16_t	signed, fastest type that is at least 16-bits
intptr_t	signed, wide enough to hold a pointer

Pointers Are 32-bit Aligned

```
_nassert((intptr_t)p1 % 2 == 0);  
_nassert((intptr_t)p2 % 2 == 0);
```

- **_nassert** is similar to an intrinsic, but generates no code
- Means the expression is always true
- Expression says the pointer is aligned to a 32-bit boundary
 - C28x addresses are counted in 16-bit words
 - Aligned to an even word boundary

Loop Runs an Even Number of Times

```
#pragma MUST_ITERATE(,,2)
```

- #pragma is preprocessor directive
- MUST_ITERATE describes behavior of the next loop
- MUST_ITERATE(min, max, multiple)
 - min: minimum number of times the loop runs
 - max: maximum number of times the loop runs
 - multiple: loop runs a multiple of this many times
- Can leave arguments blank
- (,,2) means no minimum, no maximum, multiple of 2

16x16 to 32 Multiply

```
result += (int32_t) p1[i] * p2[i];
```

- Without the cast, the upper bits of the multiply are lost
- Default behavior of multiply in C works this way
 - `(int32_t) ((int16_t) p1[i] * (int16_t) p2[i])`
 - 16x16 multiply, keep lower 16-bits, signed extend to 32-bits
- The cast changes behavior from default to this
 - `(int32_t) p1[i] * (int32_t) p2[i]`
 - Each operand is signed extended to 32-bits, then multiplied
- C28x compiler recognizes this idiom, and maps to MAC-style instructions
- Application Note: How to Write Multiplies Correctly in C Code
 - <http://www.ti.com/lit/pdf/spra683> ([link](#))

Dual MAC Required Build Options

- --opt_level=2 or higher
- --unified_memory
 - Indicates data memory bus and program memory bus are connected to the same blocks of memory
 - Very rare for these buses to be connected to different blocks of memory
 - If there is any doubt, check the data sheet

Dual MAC Final Result

```
for (i = 0; i < length; i++)  
    result += (int32_t) p1[i] * p2[i];
```

Becomes ...

	RPT	AR5
	DMAC	ACC:P, *XAR4++, *XAR7++
	ADDL	P, ACC

Dual MAC How to Align Data

- Data passed to `compute_mac` function must be aligned to 32-bits
- How is that done?
- Global or static data, use `#pragma DATA_ALIGN`

```
#define N 256
#pragma DATA_ALIGN(ar1, 2)
#pragma DATA_ALIGN(ar2, 2)
int16_t ar1[N];
int16_t ar2[N];
...

compute_mac(ar1, ar2, N);
```

- Addresses returned by RTS function `malloc` are 32-bit aligned

Agenda

- Enable Optimization
- Prefer Signed Loop Counters
- Background
- Restrict
- Dual MAC – C Code Method
- **Dual MAC – Intrinsic Method**

Dual MAC – Intrinsic Method

- Second method for implementing dual MAC
- Uses intrinsic `__dmac`
- Guarantees DMAC is emitted
- Use the method you prefer

Dual MAC – Intrinsic Method

```
#include <stdint.h>

int32_t intrinsic_mac(int16_t *p1, int16_t *p2, int_fast16_t length)
{
    int_fast16_t i;
    int32_t *p1_32 = (int32_t *) p1;
    int32_t *p2_32 = (int32_t *) p2;
    int32_t acc1, acc2;

    acc1 = acc2 = 0;
    length >>= 1;

    for (i = 0; i < length; i++)
        __dmac(p1_32[i], p2_32[i], acc1, acc2, 0);

    return acc1 + acc2;
}
```

Pointers Must Be 32-bit Aligned

```
int32_t *p1_32 = (int32_t *) p1;  
int32_t *p2_32 = (int32_t *) p2;
```

- `__dmac` requires 32-bit memory operands
- These lines presume `p1` and `p2` are aligned to 32-bits
- Because of optimization, these assignments result in no code

Two Accumulators

```
int32_t acc1, acc2;
```

- `__dmac` requires two 32-bit wide variables
- Accumulate results across multiple calls to `__dmac`
- Passed by reference
 - Concept borrowed from C++
 - Internet search on *c++ reference*
 - Means the variable is modified by `__dmac`

```
return acc1 + acc2;
```

- Final result requires adding the accumulators

Halve the Loop Count

```
length >>= 1;
```

- It is called dual MAC because it computes two multiply-accumulates at once
- Thus, run the loop half as many times
- Right-shift by 1 is the same as divide by 2
- Presumes **length** is even

References

- C2000 Performance Tips and Tricks
http://processors.wiki.ti.com/index.php/C2000_Performance_Tips_and_Tricks ([link](#))
- TMS320C28x Optimizing C/C++ Compiler User's Guide
<http://www.ti.com/lit/pdf/spru514> ([link](#))

Questions?